

Floating Point Representation

20 January 2018

BACKGROUND

The usual representation of a fractional number as

$$b_{n-1}b_{n-2}\dots b_1b_0.b_{-1}b_{-2}\dots b_m$$

may not always serve our purpose.

This form

- is not suitable to represent a very big or a very small number.
 - ▶ Consider the number 5×2^{40} and it would require the string 101 followed by 40 zeros. Clearly, a 32 bit representation would fail.
- it has the accuracy issues as well.

Though the typical notation of representing a big or a small number was known to us for long (the floating point notation; e.g., Avagadro's Number = 6.022×10^{23} or the charge of an electron = -1.602×10^{-19} C) the same was not formally used in computers until the widespread acceptance of the IEEE 754 floating point standard.

Problem: No common standard

Prior to 1980 there was no common standard followed by the manufacturer (h/w and s/w) and emphasis was more on the ease of representation rather than accuracy; this leads to

- major portability and
- accuracy issues

for scientific packages involving floating point calculations.

Fortunately, IEEE 754 is now followed universally and the issues have been solved across all computing platforms.

INTEL in the 1980s appointed Prof. Kahan of University of California at Berkeley to propose a standard to be used for their 8087 floating point coprocessor; a supporting hardware for 8086/88 CPU to improve floating point performance in the PCs'. The proposal submitted by Prof. Kahan was ultimately used by the IEEE appointed floating point standard committee with minor modification.

IEEE 754 notation

Typical floating point representations uses 32 bits (Single precision) and 64 bits (Double precision).

	Sign bit (s)	Exponent (e)	Fraction (f)
Single Precision	1-bit (b_{31})	8-bits (b_{30-23})	23-bits (b_{22-0})
Double Precision	1-bit (b_{63})	11-bits (b_{62-52})	52-bits (b_{51-0})

The numeric value of the floating point number is computed as

$$V = (-1)^s M \times 2^E \text{ where}$$

- s is the sign bit
- E is the biased exponent
- M is the significand

IEEE 754 ... continued

- A bias (b) $2^{(\text{no. of bits in } e-1)} - 1$ (127 for single precision) is subtracted from the e -bits to form E (i.e., $E = e - b$).
[Note : if $e = 0$ then $E = 1 - b$]
- The fractional part (f) is stored as a 23-bit string from the implied binary point to the left of the leftmost bit (MSbit) of f . For processing a Significand M is computed as $M = 1.0 + f$ to form a normalised significand such that $1 \leq M < 2$. This virtual extra bit is used to expand the Significand by 1 bit; free of cost.
[Note: M is taken as $0.0 + f$ when $e = 0$]

Normalised, De-normalised & NaN

Depending on the bit string in the exponent part (e) the real numbers stored using IEEE 754 are interpreted as i) Normalised numbers; the common case., ii) Denormalised number and iii) Special values; like a) Infinity b) Not a Number (NaN).

Normalised	s	$e = 1 \dots 254$	f
De-normalised	s	$e = 0000 \ 0000$	f
Infinity	s	$e = 1111 \ 1111$	$f = 000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000$
NaN	s	$e = 1111 \ 1111$	$f \neq 0$

- **Normalised numbers**

- ▶ e can be any 8-bit value except 0 and 255.
- ▶ Consequently for a bias b of 127 $E = e - b$ will have the range of -126 to 127 .

- f represents any 23 bit string and the significand is normalised by adding 1.0 with f giving $1 \leq M < 2$ range to M .

Normalised, De-normalised & NaN ... continued

Denormalised number is a requirement as through the normalised representation you

- Cannot represent zero; and
- small numbers close to zero may have less accuracy.

The e bits would be zero indicating a **De-normalised** representation for which

- $E = 1 - b$; and
- $M = 0.f$, i.e., 1.0 would not be added and M will have the range of $0 \leq M < 1$

f represents any 23 bit string for single precision.

Normalised, De-normalised & NaN ... continued

Special and NaN is a clever way of representing non-real numbers or something which are exceptional cases.

Here, e bit string consists of all 1s. And we have two different cases.

- If the f part is zero then the representation indicates $+\infty$ or $-\infty$ depending on the sign bit s .
- If the f part is non-zero then the representation indicates any NaN indicating something which is not a real number. Say, for example $\sqrt{-1}$ or a real variable not yet assigned any legitimate value or divide by zero error by using some agreed upon pattern of f bit string.

Examples of Floating Point Encoding : Normalized

As a practical example let us examine how an integer $x = 12345$ is encoded if we change it to FP (i.y., float $y = (\text{float}) x$;

The binary encoding of $12345_{10} = 11\ 0000\ 0011\ 1001$. We create a normalized representation of the above by

$$12345_{10} = 1.\mathbf{100\ 0000\ 1110\ 01} \times 2^{13}$$

[Note: The binary point is moved towards left by 13 position; so to keep the value same in normalized form we need to multiply it by 2^{13}] So, the final form in single precision is

0 10001100 **100 0000 1110 01** 0000000000

[Note: MSB is $s=0$; $e = 140$; Therefore $E = e - 127$ (the bias) $= 13$; reflecting multiplication of the normalized M by 2^{13} ; Also see that the ten 0s' in the LSB position to make f to 23-bits. Finally, the leading 1 of M is not stored]

Examples of Encoding : De-Normalized and NaN

Value	32-bit Hex (single precision)	Type
0	00 00 00 00	De-normalized
-0	00 00 00 00	De-normalized
0.0	00 00 00 00	De-normalized
-0.0	80 00 00 00	De-normalized
$\sqrt{-1}$	FF C0 00 00	NaN
$-\sqrt{-1}$	7F C0 00 00	NaN
$1/(1.609\text{e-}39)$	7F 80 00 00	Special ($+\infty$)
$-1/(1.609\text{e-}39)$	FF 80 00 00	Special ($-\infty$)

float variable x is assigned different values and the corresponding memory storage is displayed in Hex to see the encoding for comprehension.

[Ref: Computer Systems: A Programmer's Perspective by R E BRYANT and D R O'Hallaron, 3rd Edition, Pearson India]

Examples of Floating Point Encoding: Continued

Let us see the encoding of a normalized negative real number, say $x = -12345.75$. The binary encoding is $-12345.75_{10} = -11\ 0000\ 0011\ 1001.11$. Note, the presence of 11 after the binary point (.). The corresponding Floating point is given below:

1 10001100 **100 0000 1110 01** 11 00000000

[Note: MSB is $s=1$; $e = 140$; Therefore $E = e - 127$ (the bias) $= 13$; reflecting multiplication of the normalized M by 2^{13} ; Also see i) the presense of 11 (before the trailing eight zeroes that reflects) $2^{-1}(= 0.5) + 2^{-2}(= 0.25) = 0.75$ and; ii) the eight 0s' in the LSB position to make f to 23-bits. Finally, the leading 1 of M is not stored]

Range

It may be noted that the difference between a 32-bit integer representation and the corresponding floating point representation is

- the larger dynamic range for the latter;
- all the integers are equally spaced in the number line;
- the real numbers are not equally spaced in the number line;
- Most of the real numbers cannot be exactly represented (we have infinite number of real numbers between any two real numbers however close); only approximated.

Range ... contd

The following table shows some of the positive number representations in single precision.

Number	exp	frac	Value	Decimal
Zero	00 ... 00	0 ... 00	0	0.0
Smallest (Denorm.)	00 ... 00	0 ... 01	$2^{-23} \times 2^{-126}$	1.4×10^{-45}
Largest (Denorm.)	00 ... 00	1 ... 11	$(1 - \epsilon) \times 2^{-126}$	1.2×10^{38}
Smallest (Norm.)	00 ... 01	0 ... 00	1×2^{-126}	1.2×10^{-38}
One	01 ... 11	0 ... 00	1×2^0	1.0
Largest (Norm.)	11 ... 10	1 ... 11	$(2 - \epsilon) \times 2^{127}$	3.4×10^{38}

Rounding

Rounding is done to reduce inaccuracy in representing a real number. Thus the goal to represent a number x is to find out the closest match x' in a systematic manner. Instead of a close match methods to set lower (x^-) and upper bounds (x^+) such that $x^- \leq x \leq x^+$ are also used. There are 4 rounding methods used in IEEE 754 standard. The first one is to find a close match. The next three are based on setting a lower and an upper boundary.

Mode	Values to be rounded				
	1.40	1.60	1.50	2.50	-1.50
Round-to-even	1	1	1	2	-1
Round-toward-zero	1	2	2	2	-2
Round-down	1	1	1	2	-2
Round-up	2	2	2	3	-1